

MACHINE LEARNING



MODELING WITH NN

- Learning from data
- Learning as parameter estimation
- Differentiation and gradient descent
- Walking through a simple learning algorithm
- PyTorch *autograd*

MODELING WITH NN

- *Learning*

Algorithms not engineered to solve a particular problem.

- Instead, it is capable of approximating (fitting) the data using a much wider family of functions.

Key idea (supervised, deep learning):

Forward pass

Backward pass

Optimization

Repetition Forward pass

Key idea (supervised, deep learning):

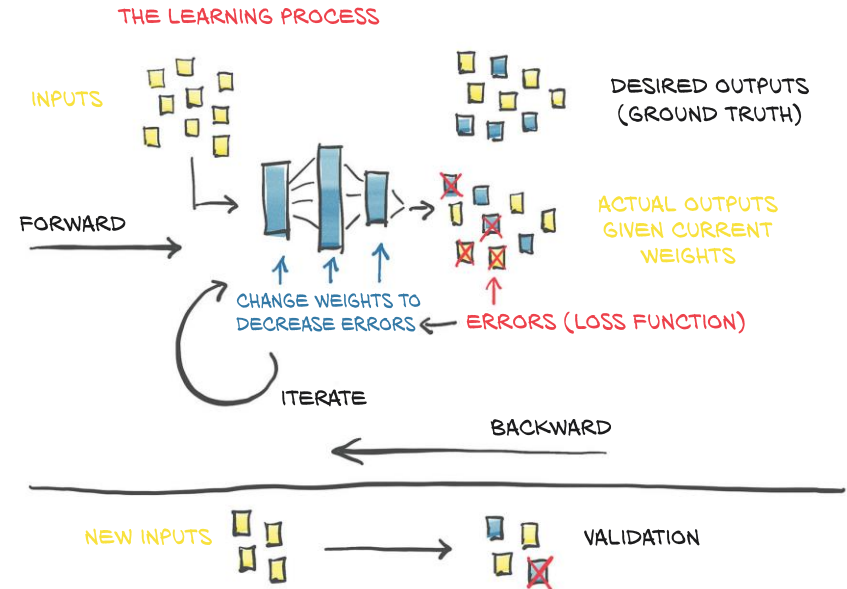
1. Forward pass

- Given input data and the corresponding desired outputs (ground truth), as well as
- initial values for the weights,
- the model is fed input data (forward pass), and a measure of the error is evaluated by comparing the resulting outputs to the ground truth.

2. Backward pass

3. Optimization

4. Repetition



Key idea (supervised, deep learning):

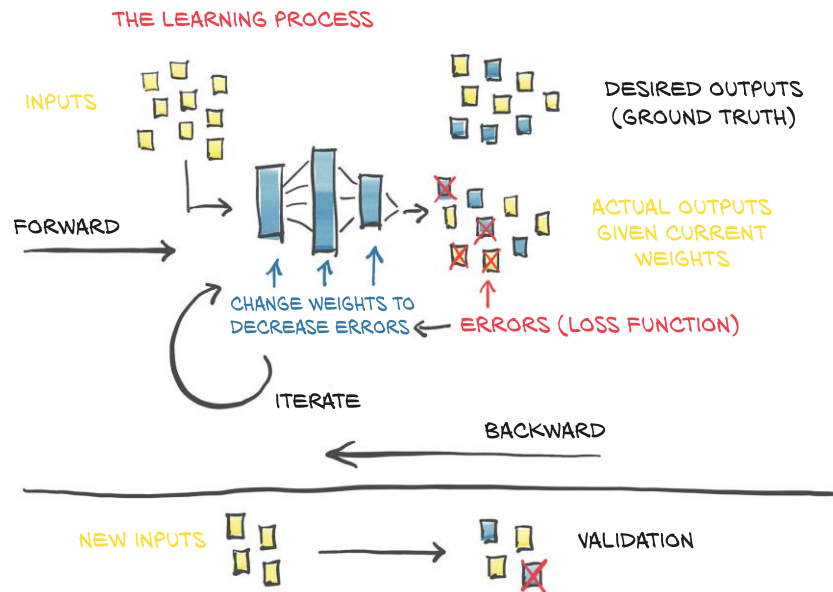
1. Forward pass

2. Backward pass

- In order to optimize the parameters of the model—its *weights and biases*—the change in the error following a unit change in weights (that is, the **gradient of the error with respect to the parameters**) is computed using the chain rule for the derivative of a composite function (backward pass).

3. Optimization

4. Repetition



Key idea (supervised, deep learning):


1. Forward pass

2. Backward pass

3. Optimization

- The value of the weights is then updated in the direction that leads to a decrease in the error.


4. Repetition



Key idea (supervised, deep learning):

1. Forward pass
2. Backward pass
3. Optimization
4. Repetition

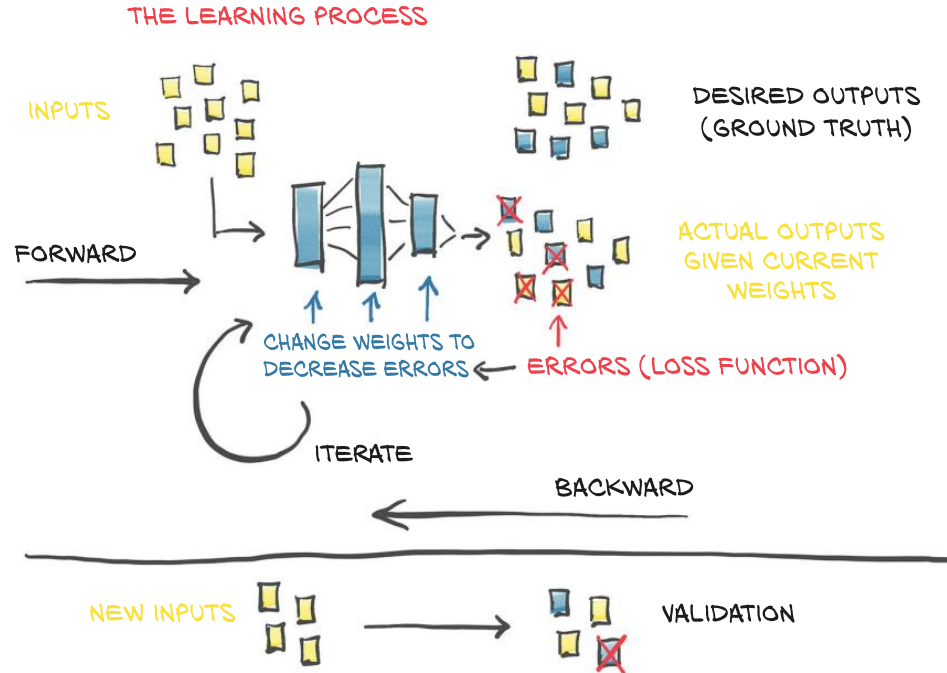
- The procedure is repeated until the error, evaluated on unseen data, falls below an acceptable level.



Key idea (supervised, deep learning)

Key concepts:


- Labeled data
- Forward pass
- Error
- Backward pass
- Update weights
- Iteration
- Validation



MODELING WITH NN

- *Learning*
- *Example: Simple Linear Regression*

Say we want to discover the equation that matches the following set of data




Pretty straightforward example

Inputs observed

[0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1,1.1,1.2,1.3,1.4,1.5,1.6,1.7,1.8,1.9,2,2.1,2.2,2.3,2.4,2.5]

Outputs observed

[0.27,0.3,0.38,0.36,0.39,0.42,0.45,0.40,0.45,0.54,0.57,0.6,0.60,0.66,0.69,0.72,0.75,0.78,0.81,0.84,0.87,0.9,0.93,0.96,0.99]





It seems to me that it's a linear function

It seems that a simple linear function should fit our observations

$$y = w x + b$$

Diagram illustrating the linear function $y = w x + b$ with labels and arrows:

- Weight** points to w (downward arrow).
- Bias** points to b (downward arrow).
- Input** points to x (upward arrow).
- Output** points to y (upward arrow).

Simplest model, linear. We want to apply this model on our set of inputs/outputs.

$m(x, w, b)$ returns y

```
def model(x, w, b):  
    y = w * x + b  
    return y
```

Training a neural network will essentially involve changing the model for a slightly more elaborate one, with a few (or millions) more parameters.

We'll go through this simple example using PyTorch.

The model has a weight w and a bias b that we want to estimate.

Our optimization process aims at finding w and b so that the loss function is at a minimum.

What is a Loss Function????

The Loss Function describes the difference between the set of values predicted by our model (\bar{y}) or y_{hat} and the set of values observed as outputs (y)

Convention in ML community

\bar{y} = values of our model we
want to fit (labels, y_{hat})

y = observed values

$$L(\bar{y}, y) = |\bar{y} - y|$$

$$L(y) = (\bar{y} - y)^2$$

MSE (mean squared error) -> widely used

```
def loss_fn(y_hat, y):  
    squared_diffs = (y_hat - y)**2  
    mse = squared_diffs.mean()  
    return mse
```

As less loss as
possible

Initialize weights and calculate predicted value y

```
w = torch.ones(())
```

```
b = torch.zeros(())
```

```
y = model(x, w, b)
```

Remember



```
def model(x, w, b):
```

```
    y = w * x + b
```

```
    return y
```

Calculate loss:

Note that this returns a tensor.

```
loss = loss_fn(y_hat, y)
```

```
loss
```

```
IPdb [2]: squared_diffs
Out [2]:
tensor([0.0035, 0.0161, 0.0209, 0.0687, 0.1088, 0.1580, 0.2164, 0.3757,
        0.4365, 0.4467, 0.5418, 0.6460, 0.8126, 0.8820, 1.0138, 1.1547,
        1.3048, 1.4641, 1.6325, 1.8101, 1.9968, 2.1928, 2.3979, 2.6121,
        2.8356])
```

```
IPdb [3]: mse
Out [3]: tensor(1.0059)
```

```
IPdb [4]:
```

MODELING WITH NN

- *Learning*
- *Example: Simple Linear Regression*
- *Gradient descent*

Reminder!: What do we want to figure out ?

We want to figure out the set of (w,b) that fits our linear equation $\bar{y} = w * x + b$ to our set of observations (labels, or y)

We will use our loss function to achieve this target, the less values got by our loss function, the better.

We will try to find the set of (w,b) that minimizes our loss function

We tackle three calculus challenges in order to reach our target

Two parameters to change when minimizing the loss function, not just one \rightarrow w and b



Partial Derivatives

Our loss function is not directly connected to our parameters



Chain Rule

We want to write an algorithm that manages to find the parameters that minimizes our values



Gradient!!

Calculus provide us a tool to find minimum values

Derivatives



$$f'(x) = \lim_{h \rightarrow 0} \frac{f(x+h) - f(x)}{h}$$



Compute a derivative



Universidad
Francisco de
Vitoria

UFV Madrid

Given:

$$y = (x - 3)^2$$

Compute its derivative: $\frac{dy}{dx}$

Given a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, and an equation

$$y = f(z^1, z^2, \dots, z^n)$$

where y and z^i for $i = 1 \dots n$ are scalars, we will write the partial derivative of the function f with respect to the variable z^i as

$$\frac{\partial f(z^1 \dots z^n)}{\partial z^i}$$

Definition of a partial derivative: this is the derivative of f with respect to parameter z^i , holding the other arguments to f (i.e., z^j for $j \neq i$) fixed. Or more precisely,

$$\frac{\partial f(z^1, \dots, z^n)}{\partial z^i} = \lim_{h \rightarrow 0} \frac{f(z^1, \dots, z^i + h, \dots, z^n) - f(z^1, \dots, z^n)}{h}$$

which can be read as “the partial derivative of y with respect to z^i , under function f , at values z^1, \dots, z^n .”

Say we have the following function:

$$L(y) = (\bar{y} - y)^2$$

But y is another function that has our parameters we want to minimize:

$$\bar{y} = w \cdot x + b$$

Then if we want to calculate the derivative of the loss function with respect to w

$$\frac{\partial L(\bar{y})}{\partial w} = \frac{\partial L}{\partial \bar{y}} \cdot \frac{\partial \bar{y}}{\partial w}$$

$$\frac{\partial L(\bar{y})}{\partial b} = \frac{\partial L}{\partial \bar{y}} \cdot \frac{\partial \bar{y}}{\partial b}$$

Given a function $f : \mathbb{R}^n \rightarrow \mathbb{R}$, and an equation

$$y = f(z^1, z^2, \dots, z^n)$$

where y and z^i for $i = 1 \dots n$ are scalars, the gradient of f , ∇f , the gradient of f at a point $p = (z^1, z^2, \dots, z^n)$

$$\nabla f(p) = \begin{bmatrix} \frac{\partial f(z^1, \dots, z^n)}{\partial z^1} \\ \vdots \\ \frac{\partial f(z^1, \dots, z^n)}{\partial z^n} \end{bmatrix}$$

Note that $\nabla f : \mathbb{R}^n \rightarrow \mathbb{R}^n$ defined at point $p = (z^1, z^2, \dots, z^n)$ and returning an n -dimensional vector

MODELING WITH NN

- *Learning*
- *Example: RTD*
- *Gradient descent*

Multi-variable function $f(\mathbf{z})$ is defined and differentiable in a neighborhood of a point \mathbf{z} , then $f(\mathbf{z})$ decreases fastest if one goes from \mathbf{z} in the direction of the negative gradient of f at \mathbf{z} , $-\nabla f(\mathbf{z})$.

Given a value \mathbf{z}_0 it follows that, for a small $\gamma \in \mathbb{R}^+$, then if we perform the update:

$$\mathbf{z}_1 = \mathbf{z}_0 - \gamma \nabla f(\mathbf{z}_0)$$

If we perform multiple updates:


$$\mathbf{z}_{n+1} = \mathbf{z}_n - \gamma \nabla f(\mathbf{z}_n)$$

We update a monotonically decreasing function:

$$f(\mathbf{z}_0) \geq f(\mathbf{z}_1) \geq f(\mathbf{z}_2) \geq f(\mathbf{z}_3) \dots$$

...hopefully, the sequence converges to the desired local minimum.

$$\mathbf{z}_1 = \mathbf{z}_0 - \gamma \nabla f(\mathbf{z}_0) \text{ then } f(\mathbf{z}_1) \leq f(\mathbf{z}_0).$$



Lines are level sets (function takes constant value)

We'll optimize the loss function with respect to the parameters using the *gradient descent* algorithm.

Compute the rate of change of the loss with respect to each parameter; and

Modify each parameter in the direction of decreasing loss (we can do it because we understand what a gradient represents)

Guys!, We've got all the tools!



$$\frac{\partial L(\bar{y})}{\partial w} = \frac{\partial L}{\partial \bar{y}} \cdot \frac{\partial \bar{y}}{\partial w} \quad \left\{ \begin{array}{l} \frac{\partial L}{\partial \bar{y}} = \frac{\partial (\bar{y} - y)^2}{\partial \bar{y}} = \frac{\partial (\bar{y}^2 - 2y\bar{y} + y^2)}{\partial \bar{y}} = 2\bar{y} - 2y = 2(\bar{y} - y) \\ \frac{\partial \bar{y}}{\partial w} = \frac{\partial (w \cdot x + b)}{\partial w} = \frac{\partial (w \cdot x) + \partial (b)}{\partial w} = x \end{array} \right.$$

$$\frac{\partial L(y)}{\partial w} = \frac{\partial L}{\partial y} \cdot \frac{\partial y}{\partial w} = 2(\bar{y} - y) \cdot x$$

$$\frac{\partial L(\bar{y})}{\partial b} = \frac{\partial L}{\partial \bar{y}} \cdot \frac{\partial \bar{y}}{\partial b} \quad \left\{ \begin{array}{l} \frac{\partial L}{\partial \bar{y}} = \frac{\partial (\bar{y} - y)^2}{\partial \bar{y}} = \frac{\partial (\bar{y}^2 - 2y\bar{y} + y^2)}{\partial \bar{y}} = 2\bar{y} - 2y = 2(\bar{y} - y) \\ \frac{\partial y}{\partial b} = \frac{\partial (w \cdot x + b)}{\partial b} = \frac{\partial (w \cdot x) + \partial (b)}{\partial b} = 1 \end{array} \right.$$

$$\frac{\partial L(\bar{y})}{\partial w} = \frac{\partial L}{\partial \bar{y}} \cdot \frac{\partial \bar{y}}{\partial w} = 2(\bar{y} - y)$$

The gradient vector has two components (2-dimensional vector)

$$\nabla L(\mathbf{y}) = [2(\bar{y} - y)] \cdot x, 2(\bar{y} - y)]$$



Code your own functions

Code in Python the following functions:

- A function that, given a weight w and a bias b , returns the function $y = ax + b$
- Write a function that, given a function $y = wx + b$, returns its gradient vector

$$\frac{\partial L}{\partial y} = 2(\bar{y} - y)$$

$$\frac{\partial \bar{y}}{\partial w} = x$$

$$\frac{\partial \bar{y}}{\partial b} = 1$$

```
def dloss_fn(y_hat, y):  
    dLoss_y = 2 * (y_hat - y)  
    return dLoss_y
```

```
def dmodel_dw(x, w, b):  
    return x
```

```
def dmodel_db(x, w, b):  
    return 1.0
```

```
def grad_fn(x, y, y_hat, w, b):  
    dloss_dtp = dloss_fn(y_hat, y)  
    dloss_dw = dloss_dtp * dmodel_dw(x, w, b)  
    dloss_db = dloss_dtp * dmodel_db(x, w, b)  
    return torch.stack([dloss_dw.sum() / y_hat.size(0), dloss_db.sum() / y_hat.size(0)])
```

$$\nabla L(y) = [2(\bar{y} - y)] \cdot x, 2(\bar{y} - y)]$$

```
def grad_fn(x, y, y_hat, w, b):  
    dloss_dtp = dloss_fn(y_hat, y)  
    dloss_dw = dloss_dtp * dmodel_dw(x, w, b)  
    dloss_db = dloss_dtp * dmodel_db(x, w, b)  
    return torch.stack([dloss_dw.sum() / y_hat.size(0), dloss_db.sum() / y_hat.size(0)])
```

Ey Ey Ey Ey... Why do you include `sum()` and divide by `my_model.size()` ??

```
def grad_fn(x, y, y_hat, w, b):  
    dloss_dtp = dloss_fn(y_hat, y)  
    dloss_dw = dloss_dtp * dmodel_dw(x, w, b)  
    dloss_db = dloss_dtp * dmodel_db(x, w, b)  
    return torch.stack([dloss_dw.sum() / y_hat.size(0), dloss_db.sum() / y_hat.size(0)])
```

Reminder: we have this observations

Inputs observed

[0.1,0.2,0.3,0.4,0.5,0.6,0.7,0.8,0.9,1,1.1,1.2,1.3,1.4,1.5,1.6,1.7,1.8,1.9,2,2.1,2.2,2.3,2.4,2.5]

Outputs observed

[0.27,0.3,0.38,0.36,0.39,0.42,0.45,0.40,0.45,0.54,0.57,0.6,0.60,0.66,0.69,0.72,0.75,0.78,0.81,0.84,0.87,0.9,0.93,0.96,0.99]


Therefore, what we are going to calculate is the mean of the gradient in the points we have observed!

$$\nabla L(\mathbf{y}) = \left[\frac{\sum_{k=1}^N [(\bar{y}_k - y_k) \cdot x_k]}{N}, \frac{\sum_{k=1}^N (\bar{y}_k - y_k)}{N} \right]$$

N = number of observations

Given a set of values w_0, b_0 it follows that, for a small $\gamma \in \mathbb{R}^+$, then if we perform the update:

$$(w_1, b_1) = (w_0, b_0) - \gamma \cdot \nabla y(w_0, b_0)$$



```
w, b = params
y = model(x_hat, w, b)

gradient_vector = grad_fn(x_hat, y_hat, y, w, b)
params = params - learning_rate * gradient_vector

loss = loss_fn(y, y_hat)
losses.append(loss)
```

$\gamma = \text{learning rate}$



<https://towardsdatascience.com/the-gradient-vector-66ad563ab55a>




Roman Paolucci

Jun 5, 2020 · 4 min read · Member-only · Listen



The Gradient Vector

What is it, and how do we compute it?



98



MODELING WITH NN

- *Learning*
- *Example: RTD*
- *Gradient descent*
- ***A training loop***

All in place to start optimizing parameters.

Iterate a fixed number of iterations, or until paramters don't change (enough).

Epoch: an interation over the dataset.

```
def training_loop(n_epochs, learning_rate, params, x, y, print_params=True):  
    for epoch in range(1, n_epochs + 1):  
        # Params we want to fit  
        w, b = params  
  
        # Setting the model we will use (a simple linear equation)  
        y_hat = model(x, w, b)  
  
        # Setting our gradient vector  
        gradient_vector = grad_fn(x, y, y_hat, w, b)  
  
        # Setting our Gradient Descent step  
        params = params - learning_rate * gradient_vector  
  
        loss = loss_fn(y_hat, y)  
        losses.append(loss)
```

1. Setting our params


2. Setting our model

3. Setting our Gradient Vector

4. Gradient Descent step

Let's run our algorithm


```
Epoch 1, Loss 0.020676
  Params: tensor([0.3936, 0.2293])
  Grad:  tensor([0.4222, 0.2474])
  Loss:  tensor(0.0207)
Epoch 2, Loss 0.018354
  Params: tensor([0.3896, 0.2270])
  Grad:  tensor([0.3971, 0.2315])
  Loss:  tensor(0.0184)
Epoch 3, Loss 0.016306
  Params: tensor([0.3859, 0.2248])
  Grad:  tensor([0.3735, 0.2165])
  Loss:  tensor(0.0163)
...
Epoch 10, Loss 0.007339
  Params: tensor([0.3652, 0.2131])
  Grad:  tensor([0.2441, 0.1345])
  Loss:  tensor(0.0073)
Epoch 11, Loss 0.006586
  Params: tensor([0.3629, 0.2119])
  Grad:  tensor([0.2298, 0.1254])
  Loss:  tensor(0.0066)
...
Epoch 99, Loss 0.000765
  Params: tensor([0.3218, 0.2016])
  Grad:  tensor([ 0.0079, -0.0105])
  Loss:  tensor(0.0008)
Epoch 100, Loss 0.000763
  Params: tensor([0.3217, 0.2017])
  Grad:  tensor([ 0.0079, -0.0105])
  Loss:  tensor(0.0008)
...
Epoch 1000, Loss 0.000513
  Params: tensor([0.3020, 0.2323])
  Grad:  tensor([ 0.0003, -0.0005])
  Loss:  tensor(0.0005)
  Params: tensor([0.3020, 0.2323])
```



Check out the results

```
Our model estimates a  $w = 0.291$  and  $b = 0.249$   
Our model predicts for the input 2.800 the output 1.065
```

It seems that our algorithm has found
a model that fit our observations





Trying fit other linear regressions

- Try another set of observations
- Load your observations into the model
- Did the algorithm manage to converge?

What have we achieved

Just a new way of computing parameters of a linear regression... but:

This process can be applied to any model that depends on parameters, as long as all functions can be differentiated analytically.

The only issue is that if we change the model, we need to recompute the gradient of the loss with respect to the parameters applying the chain rule.

... this is where *autograd* and *backpropagation* comes in.



What have we achieved



Universidad
Francisco de
Vitoria

UFV Madrid



Until now, tensors were just multidimensional arrays (plus device).

But tensors objects are designed to be the building block of **computational graphs**, and to **support automatic differentiation**.

Nodes of computational graphs, connected to other nodes through functions (functional module instances)

Three key attributes:

- data
- grad
- grad_fn