

MACHINE LEARNING



BUILDING OUR FIRST NEURAL NETWORK

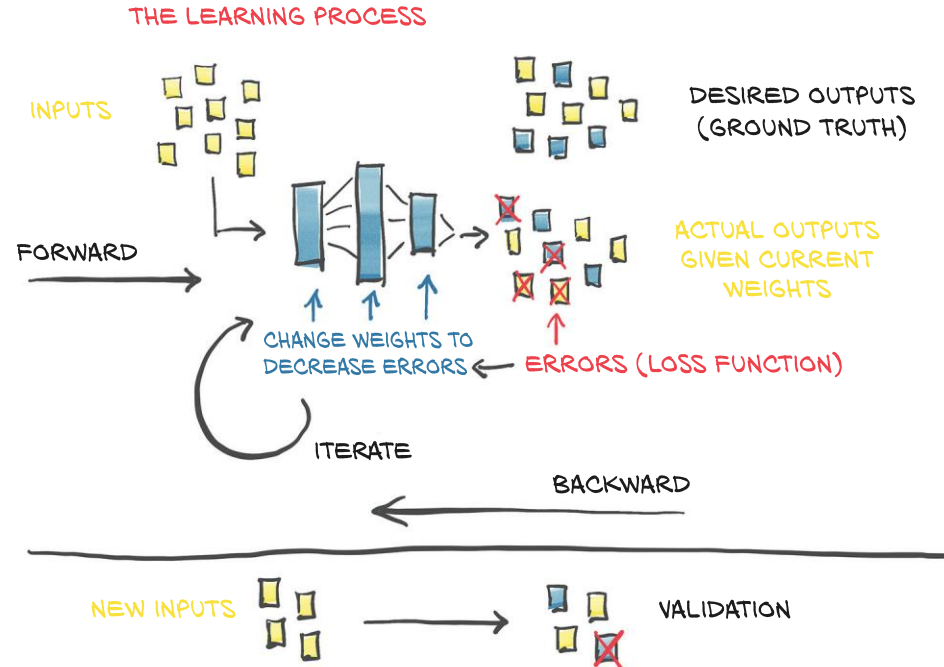
Reminder, Forward step and Backward Step

Forward pass


Backward pass

Optimization

Repetition Forward pass



We will move from a Linear Regression to Neural Network



We will go through this evolution using a binary classification problem

Problem: to classify pictures as “picture” or “Non-Cat” picture



1 (cat)

0 (cat)

y (label)

How is a picture represented in a computer ?

				122	13	64	230
				252	123	34	23
255	231	30	20				
123	94	83	3				
34	76	67	32				
12	23	11	201				

RGB (Red, Green, Blue)

64 pixels




64 pixels



3 Arrays (64x64)

X
(feature)

We will format our 3-dimensional array into a single vector x




x Vector with $64 \cdot 64 \cdot 3 = 12288$ components

y Vector with just 1 component

Notation

Given a sample $(x, y) \rightarrow x \in \mathbb{R}^{n_x}, y \in \{0, 1\}$


m training samples: $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$



A logistic regression is a Neural Network used for binary classification

Our problem: Given a sample x , we want to \hat{y} as the probability of being a cat $\hat{y} = P(y = 1|x) \quad 0 \leq \hat{y} \leq 1$

Our model: $z = w^T \cdot x + b$, where $w \in \mathbb{R}^{n_x}$, $b \in \mathbb{R}$




But... , this function returns values greater than 1 !!!



And.., probability should be a real number between, and 1

Here it comes... the neuron!



We call $z = w^T + b$

$$\hat{y} = \sigma(z) \quad \sigma(z) = \frac{1}{1 + e^{-z}}$$

If z is large + ...

$$\sigma(z) \approx \frac{1}{1 + 0} \approx 1$$

If z is large -...

$$\sigma(z) \approx \frac{1}{1 + \text{big}} \approx 0$$

Now our output is
number between
0 and 1

What is a neuron ?


A neuron in Machine Learning is a function composition between:

- linear equation (regression function)
- and a non-linear equation (activation function)

$$\hat{y} = g \circ f(w, b) = g(f(w, b))$$

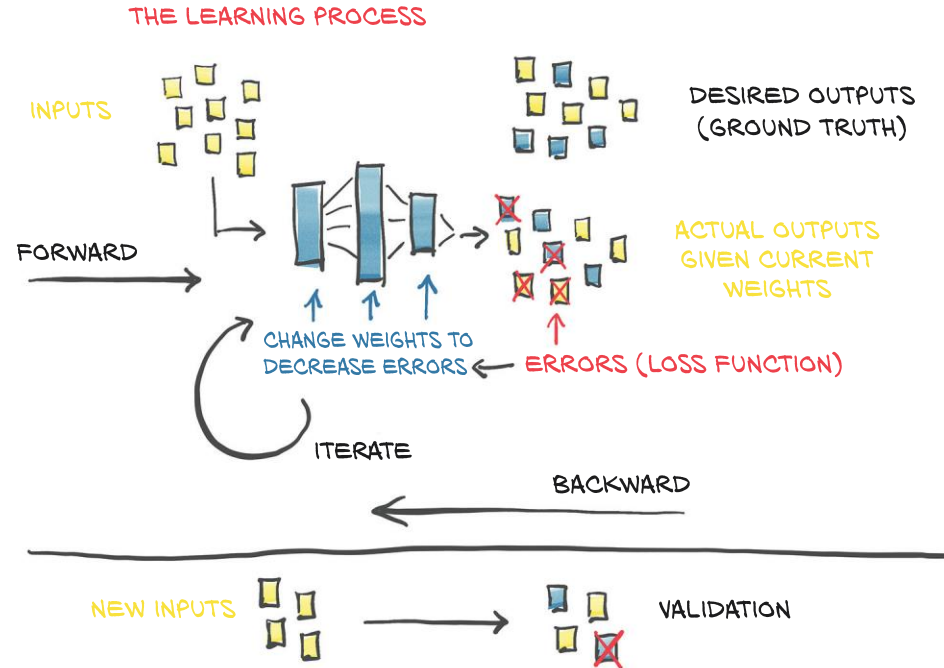
In our Cat classifier (binary classification):

$$\hat{y} = \sigma \circ z(w, b) = \sigma(z(w, b)) = \frac{1}{1 + e^{-z}} = \frac{1}{1 + e^{-(w^T x + b)}}$$



w^T is vector, b is a real number

- ✓ Forward pass
- ➡ Backward pass
- Optimization
- Repetition Forward pass



Logistic Regression Cost Function


Given $\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$,

i is the i-th example




we want $\hat{y}^{(i)} \approx y^{(i)}$

Our cost function should be $\mathcal{L}(\hat{y}, y) = (\hat{y} - y)^2$



Local
minimum
issue



Gradient Descent doesn't work well with
local minimum

Specific Loss Function for Logistic Regression

Loss function $\rightarrow \mathcal{L} = (\hat{y}, y) - (y \cdot \log(\hat{y}) - (1 - y) \cdot \log(1 - \hat{y}))$

If $y = 1, \mathcal{L}(\hat{y}, y) = -\log(\hat{y})$ If \hat{y} close to 1, $\mathcal{L} \approx 0$

If $y = 0, \mathcal{L}(\hat{y}, y) = -\log(1 - \hat{y})$ If \hat{y} close to 0, $\mathcal{L} \approx 0$

$$\mathfrak{J}(w, b) = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)}) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})]$$


Cost function

Loss Function versus Cost Function

Loss function measures de accuracy in a single example

$$\mathcal{L}(\hat{y}, y) = -(y \cdot \log(\hat{y}) - (1 - y) \cdot \log(1 - \hat{y}))$$

Cost function measures de accuracy on the whole dataset

$$\mathfrak{J}(w, b) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log \hat{y}^{(i)} + (1 - y^{(i)}) \log(1 - \hat{y}^{(i)})]$$

We're going to suppose that we only have 1 sample with two features

x_1

w_1

x_2

w_2

b

$$z = w_1 \cdot x_1 + w_2 \cdot x_2 + b \Rightarrow a = \sigma(z) \Rightarrow \mathcal{L}(a, y)$$

First
backward step

Second
backward step

Third
backward step

$$\frac{d\mathcal{L}(a, y)}{da}$$

$$\frac{d\mathcal{L}(a, y)}{dz}$$

$$\left[\frac{\partial \mathcal{L}(a, y)}{\partial w_1}, \frac{\partial \mathcal{L}(a, y)}{\partial w_2}, \frac{\partial \mathcal{L}(a, y)}{\partial b} \right]$$

Ground
truth

output

feature

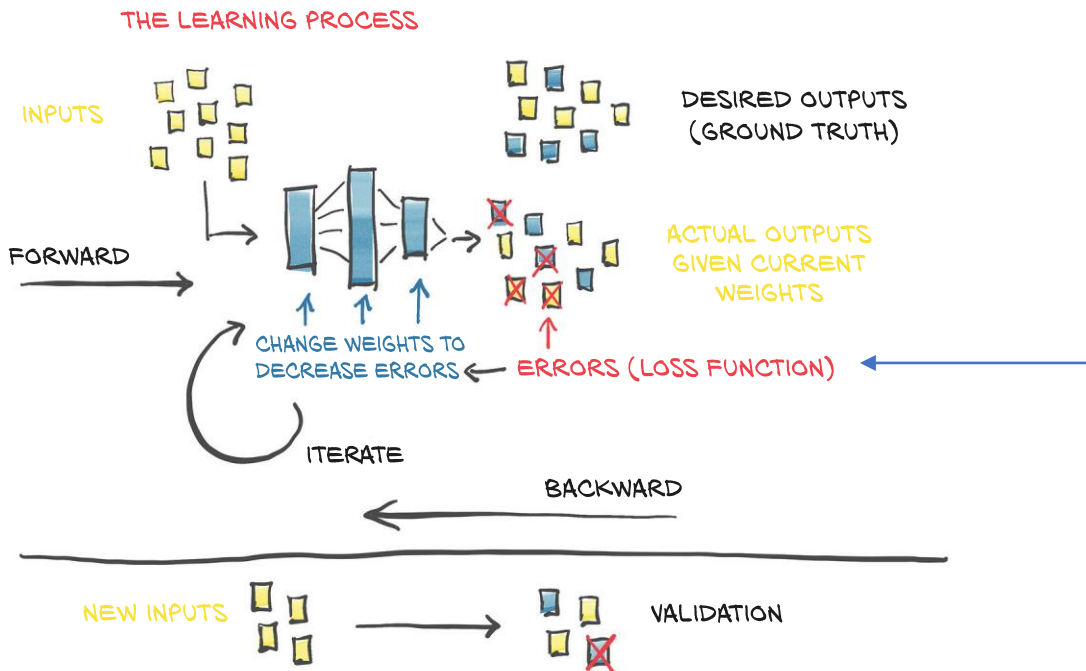
$$\frac{\partial \mathcal{L}(a, y)}{\partial w_1} = \frac{dz}{dw_1} \cdot \frac{da}{dz} \cdot \frac{d\mathcal{L}(a, y)}{da} = x_1 \cdot a \cdot (1 - a) \cdot \left(-\frac{y}{a} + \frac{1 - y}{1 - a} \right) = x_1 (a - y)$$

Therefore, the gradient vector for the cost function is

$$\nabla \mathcal{L}[w_1, w_2, b] = [x_1(a - y), x_2(a - y), (a - y)]$$

- ✓ Forward pass
- ✓ Backward pass
- ➡ Optimization

Repetition Forward pass



Remind the gradient descent algorithm used in the linear regression practice


$$\mathbf{z}_{n+1} = \mathbf{z}_n - \gamma \nabla \mathcal{L}(\mathbf{z}_n)$$

$$w_{1_{n+1}} = w_{1_n} - \gamma \nabla \mathcal{L}(w_{1_n}) = w_{1_n} - \gamma \cdot x_n (a_n - y_n)$$

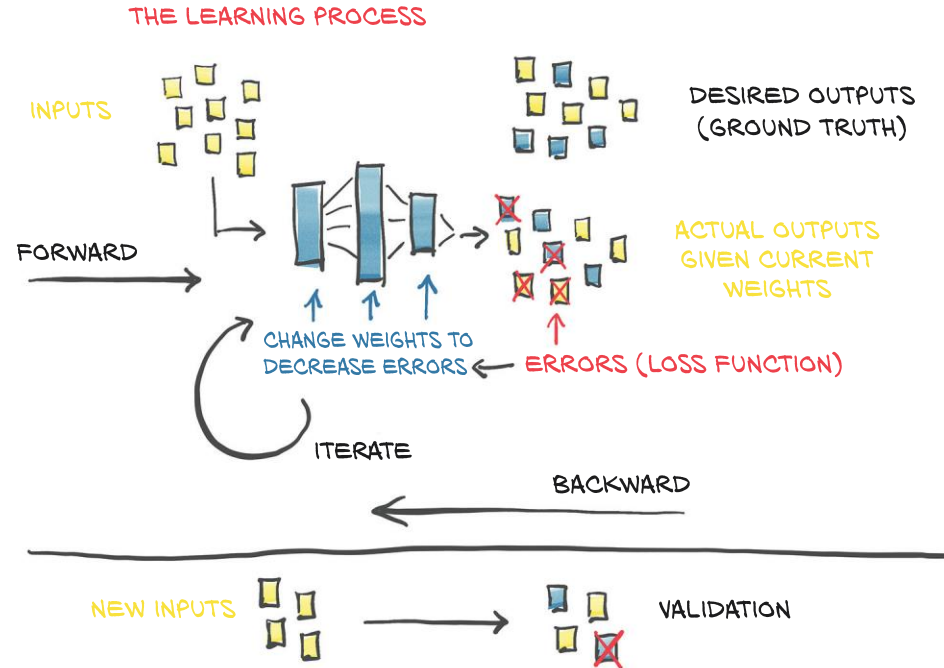
$$w_{2_{n+1}} = w_{2_n} - \gamma \nabla \mathcal{L}(w_{2_n}) = w_{2_n} - \gamma \cdot x_n (a_n - y_n)$$

$$b_{n+1} = b_n - \gamma \nabla \mathcal{L}(b_n) = b_n - \gamma \cdot (a_n - y_n)$$

γ = Learning rate



- ✓ Forward pass
- ✓ Backward pass
- ✓ Optimization
- ➡ Repetition Forward pass




We will use a for loop to iterate as many times as we consider (epochs)

```
for epochs in range(1,n_epochs+1)

    a = compute our output
    Loss = compute our loss
    Gradient = compute our gradient
    Params in epoch n = (params in epoch n-1)- learning rate x Gradient Descent
```

- ✓ Forward pass
- ✓ Backward pass
- ✓ Optimization
- ✓ Repetition Forward pass



How will our algorithm be with m examples instead of just
1 example?

Logistic regression with m examples instead of 1 example

What happens if I have to compute the algorithm for m examples instead of 1 example??

Let's review our trainin loop

```
for epoch in epochs:
```

```
    for i in range(1,m+1):
```

$$z^{(i)} = w^T \cdot x^{(i)} + b$$

$$a^{(i)} = \sigma(z^{(i)})$$

$$\mathcal{L} = \mathcal{L} + (-[y^{(i)} - \log a^{(i)} + (1 - y^{(i)}) \log(1 - a^{(i)})])$$

$$dz^{(i)} = a^{(i)} - y^{(i)}$$

$$db = db + dz^{(i)}$$

```
    for n in range(1,n_features):
```

$$dw_n = x_n^{(i)} \cdot dz^{(i)}$$

$$\mathfrak{J} = \mathcal{L}/m, dw_1/m, \dots, dw_n/m, db/m$$

```
    for n in range(1,n_features):
```

$$w_n = w_n - \gamma \cdot dw_n$$

Forward Function

Loss Function

Gradient Vector

Cost Function

Optimization

$o(n^3)$



MACHINE LEARNING



VECTORIZATION

**$o(n^3)$ algorithms don't escalate, moreover in Machine Learning environments,
where data tends to grow exponentially**

Non-Vectorized

$$z = w^T \cdot x + b$$

$$z = 0$$

for i in range (1,n_features+1)

$$z = z + w[i] \cdot x[i]$$

Vectorized

$$z = np.dot(w, x)$$

$$w^T x$$

Performance Comparasion

```
# -*- coding: utf-8 -*-
"""
Created on Thu Oct 20 10:59:15 2022

@author: PedroAnquela
"""

import numpy as np
import time

dim = 10000000

w = np.random.rand(dim)
x = np.random.rand(dim)

start = time.time()

a = np.dot(w,x)

end = time.time()

print("Vectorized version: {0:.4f} ms.".format(1000*(end-start)))

#####
# Non-Vectorized Version #
#####
c = 0
start = time.time()

for i in range (dim):
    c += w[i] * x[i]

end = time.time()

print("Non-Vectorized version: {0:.4f} ms.".format(1000*(end-start)))
```

```
In [8]: runfile('C:/Users/PedroAnquela/Documents/Proyectos_Local/UFV/
prog_ii_2020/ml/vectorization/vectorization.py', wdir='C:/Users/
PedroAnquela/Documents/Proyectos_Local/UFV/prog_ii_2020/ml/
vectorization')
Vectorized version: 9.9742 ms.
Non-Vectorized version: 3460.2816 ms.
```

**384 times
faster
when vectorizing**

We turn data into numpy vectors

$$A = [a^{(1)}, \dots a^{(m)}]$$

$$Y = [y^{(1)}, \dots y^{(m)}]$$

$$dZ = A - Y$$

$$db = \frac{1}{m} \sum_{i=1}^m dz^{(i)} = \frac{1}{m} \text{np.sum}(dZ)$$

$$dw = \frac{1}{m} X \cdot dz^T = \frac{1}{m} \begin{bmatrix} x^{(1)} & \dots & x^{(m)} \end{bmatrix} \begin{bmatrix} dz^{(1)} \\ \dots \\ dz^{(m)} \end{bmatrix}$$

$$(1, m) \cdot (m, 1)$$

Let's review our trainin loop

$$J = 0, dw_1 = 0, dw_2 = 0, db = 0$$

for epoch in epochs:

for i in range(1,m+1):

$$z^{(i)} = w^T \cdot x^{(i)} + b$$

$$a^{(i)} = \sigma(z^{(i)})$$

$$\mathcal{L} = \mathcal{L} + (-[y^{(i)} - \log a^{(i)} + (1 - y^{(i)})\log(1 - a^{(i)})])$$

$$dz^{(i)} = a^{(i)} - y^{(i)}$$

$$db = db + dz^{(i)}$$

for n in range(1,n_features):

$$dw_n = x_n^{(i)} \cdot dz^{(i)}$$

$$\mathfrak{J} = \mathcal{L}/m, dw_1/m, \dots, dw_n/m, db/m$$

for n in range(1,n_features):

$$w_n = w_n - \gamma \cdot dw_n$$

for epoch in epochs:

$$Z = w^T X + b = np.dot(w.T, X) + b$$

$$A = \sigma(Z)$$

$$dZ = A - Y$$

$$db = \frac{1}{m} \sum_{i=1}^m \frac{1}{m} np.sum(dZ)$$

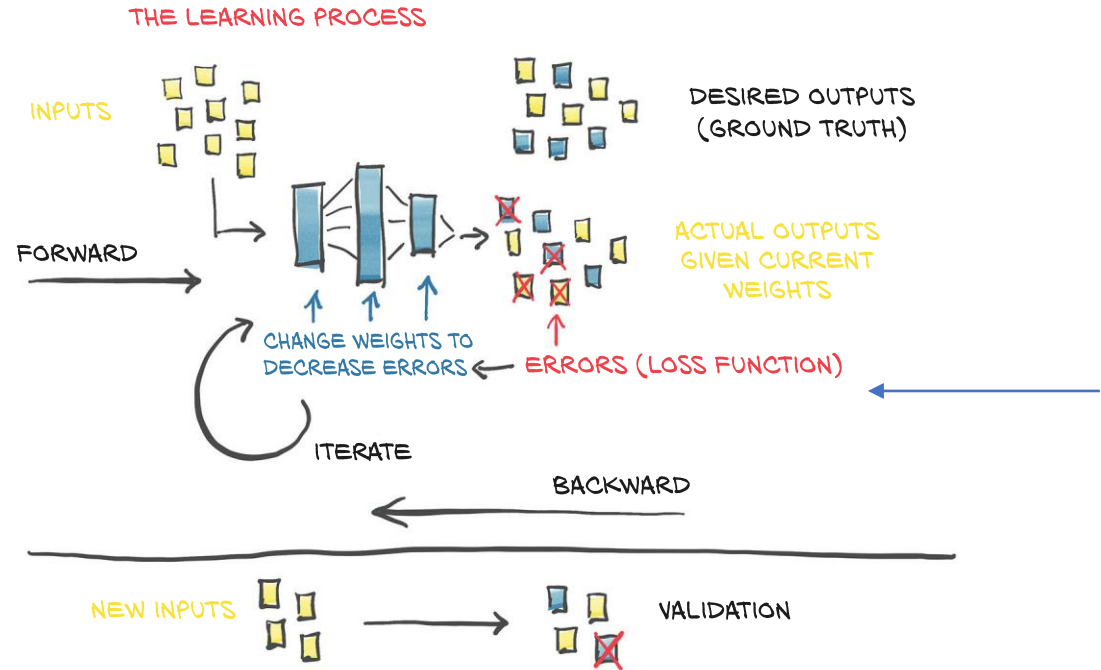
$$dw = \frac{1}{m} X \cdot dz^T$$

$$w = w - \gamma dw$$

$$b = b - \gamma db$$

Model completed!

- ✓ Forward pass
- ✓ Backward pass
- ✓ Optimization
- ✓ Repetition Forward pass



Remembering a basic concept of Linear Algebra

$$\begin{bmatrix} 100 \\ 101 \\ 102 \\ 103 \end{bmatrix} + 100 =$$


Can I do this operation ?

Two matrices must have an equal number of rows and columns to be added.

In which case, the sum of two matrices **A** and **B** will be a matrix which has the same number of rows and columns as **A** and **B**.

The sum of **A** and **B**, denoted **A + B**, is computed by adding corresponding elements of **A** and **B**.

Source: https://en.wikipedia.org/wiki/Matrix_addition





Quick snippet

- Open your text editor
- Generate a numpy array with the following dimensions (4,1)
- add the number 100 to your numpy.array.
- Print the result
- Did it work? What was the result?

```
import numpy as np

a = np.array([[100],
              [101],
              [102],
              [103]
              ])

print(a + 100)
```

Numpy broadcasts non-adjusted shapes to allow the addition

$$\begin{bmatrix} 100 \\ 101 \\ 102 \\ 103 \end{bmatrix} + 100 = \text{numpy broadcasting} = \begin{bmatrix} 100 \\ 101 \\ 102 \\ 103 \end{bmatrix} + \begin{bmatrix} 100 \\ 100 \\ 100 \\ 100 \end{bmatrix} = \begin{bmatrix} 200 \\ 201 \\ 202 \\ 203 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 2 & 5 \\ 3 & 4 & 6 \end{bmatrix} + [100 \quad 200 \quad 300] = \text{numpy broadcasting} = \begin{bmatrix} 1 & 2 & 5 \\ 3 & 4 & 6 \end{bmatrix} + \begin{bmatrix} 100 & 200 & 300 \\ 100 & 200 & 300 \end{bmatrix} = \begin{bmatrix} 101 & 202 & 305 \\ 103 & 204 & 306 \end{bmatrix}$$

Numpy broadcasts non-adjusted shapes to allow the addition


Numpy broadcasts general principles

$$A_{(m,n)} + B_{(1,n)} \longrightarrow A_{(m,n)} + B_{(m,n)}$$

$$A_{(m,n)} - C_{(m,1)} \longrightarrow A_{(m,n)} - C_{(m,n)}$$

$$A_{(m,n)} + n \in \mathbb{R} \longrightarrow A_{(m,n)} + B_{(m,n)}, b_{ij} = n \forall b_{ij} \in B$$

Never let Numpy broadcast work in Matrix Multiplication

$$A_{(m,n)} \cdot B_{(n,m)}$$


Number of columns of first matrix HAS TO match the number of rows of the second matrix



Matrix multiplication



Universidad
Francisco de
Vitoria

UFV Madrid

- Given the following matrix

$$A = \begin{bmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \end{bmatrix}$$

- And the following vector

$$B = [2, 3, 4]$$

- Use numpy to calculate $A \cdot B$

Use the command `numpy.reshape` to fix the dimensions you need.

Never let numpy decide by itself dimensions of your matrix

```
import numpy as np
```

```
a = np.ones(5)
```

```
print(a)
```

```
print(a.shape)
```

```
a = a.reshape(5,1)
```

```
print(a)
```

```
print(a.shape)
```

```
a = a.reshape(1,5)
```

```
print(a)
```

```
print(a.shape)
```

```
[1. 1. 1. 1. 1.]  
(5,)
```

```
[[1.]  
 [1.]  
 [1.]  
 [1.]  
 [1.]]  
(5, 1)
```

```
[[1. 1. 1. 1. 1.]]  
(1, 5)
```



Cat / non-Cat classifier

- Build a Logistic Regression (Shallow Neural Network) to recognize cats
- Classify your own photo as Cat or Non-Cat
- Loops are not permitted, unless there is an explicit function to do so
- Use Jupyter Notebook to fill up the Assignment (conda install jupyter)
- Clone the Notebook from the following link
- At the top of a Given Block you will find a #Graded Función tag, you will write your own code between the `###Start Solution Here ###` and the `###End Solution Here###` comments
- Take your time, think about all the concepts explained in the previous lectures. Check out if your model has converged and it can recognize a cat