

MACHINE LEARNING



CONVOLUTIONS


Understanding convolution

Creating custom nn.Module subclasses


Building a convolutional neural network

Supervised Learning


1. Feed Forward Neural Networks




2. Convolutional Neural Networks



3. Recurrent Neural Networks



4. Encoder-Decoder Architectures




Unsupervised Learning

5. Autoencoder




6. Generative Adversarial Networks



Reinforcement Learning

7. Networks for Learning Actions, Values, and Policies





CONVOLUTIONS

- *Understanding convolutions*

With the nn.Linear design, we have lost the spatial structure of the data:

- Consider the effect of the translation of an image of an airplane in the classifier
- Consider the relevance of neighboring pixels when interpreting a pixel as forming part of wing



$x =$

255
231
30
201
252
123
201

No matter what the spatial
distribution is, the x_{train} vector
will be the same

Objective, to create:

- a translation invariant model, that
- retains spatial relationship between pixels (locality)

Convolution, or more precisely, discrete is defined for a 2D image as the scalar product of a weight matrix, the kernel, with every neighborhood in the input.

The Kernel

Kernel/Filter K: the element involved in carrying out the convolution operation in the first part of a Convolutional Layer

We select K as a 3x3x1 matrix.



Kernel/Filter, K =

| | | |
|---|---|---|
| 1 | 0 | 1 |
| 0 | 1 | 0 |
| 1 | 0 | 1 |

| | | | | |
|------------------------|------------------------|------------------------|---|---|
| 1 <small>x1</small> | 1 <small>x0</small> | 1 <small>x1</small> | 0 | 0 |
| 0 <small>x0</small> | 1 <small>x1</small> | 1 <small>x0</small> | 1 | 0 |
| 0 <small>x1</small> | 0 <small>x0</small> | 1 <small>x1</small> | 1 | 1 |
| 0 | 0 | 1 | 1 | 0 |
| 0 | 1 | 1 | 0 | 0 |

Image

| | | |
|---|--|--|
| 4 | | |
| | | |
| | | |

Convolved
Feature

**So, why is this Convolution
useful for our purpose?**

$$\begin{bmatrix} 1 & 1 & 1 \\ 0 & 1 & 1 \\ 0 & 0 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix} =$$

$$1 \cdot 1 + 0 \cdot 0 + 0 \cdot 1 + 1 \cdot 0 + 1 \cdot 1 + 0 \cdot 0 + 1 \cdot 1 + 1 \cdot 0 + 1 \cdot 1 = 4$$

$$\begin{bmatrix} 1 & 1 & 0 \\ 1 & 1 & 1 \\ 0 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix} =$$

$$1 \cdot 1 + 1 \cdot 0 + 0 \cdot 1 + 1 \cdot 0 + 1 \cdot 1 + 1 \cdot 0 + 0 \cdot 1 + 1 \cdot 0 + 1 \cdot 1 = 3$$

$$\begin{bmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 1 & 1 & 1 \end{bmatrix} * \begin{bmatrix} 1 & 0 & 1 \\ 0 & 1 & 0 \\ 1 & 0 & 1 \end{bmatrix} =$$


$$1 \cdot 1 + 1 \cdot 0 + 1 \cdot 1 + 0 \cdot 0 + 1 \cdot 1 + 1 \cdot 0 + 0 \cdot 1 + 0 \cdot 0 + 1 \cdot 1 = 4$$

The convolution operators helps the model to define the Edge of the figures


Edge Detection

$$\begin{bmatrix} 10 & 10 & 10 & 0 & 0 & 0 \\ 10 & 10 & 10 & 0 & 0 & 0 \\ 10 & 10 & 10 & 0 & 0 & 0 \\ 10 & 10 & 10 & 0 & 0 & 0 \\ 10 & 10 & 10 & 0 & 0 & 0 \\ 10 & 10 & 10 & 0 & 0 & 0 \end{bmatrix} * \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix} = \begin{bmatrix} 0 & 30 & 30 & 0 \\ 0 & 30 & 30 & 0 \\ 0 & 30 & 30 & 0 \\ 0 & 30 & 30 & 0 \\ 0 & 30 & 30 & 0 \\ 0 & 30 & 30 & 0 \end{bmatrix}$$

It has found
the edge



Edge detection is the way used by the living beings to recognize the environments



Vertical Edge



Horizontal Edge

**The classic literature of convolutions is plenty of different Kernels for
Edge Detection**

$$\begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix}$$

Vertical Edge

$$\begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$

Horizontal Edge

$$\begin{bmatrix} 1 & 0 & -1 \\ 2 & 0 & -2 \\ 1 & 0 & -1 \end{bmatrix}$$

Sobel Filter

$$\begin{bmatrix} 3 & 0 & -3 \\ 10 & 0 & -10 \\ 3 & 0 & -3 \end{bmatrix}$$

Scharr Filter

**How do we choose the best filter for our convolution (for example
cats?**

We've got a powerful tool now!

Let's the Machine Learning choose the kernel for us!

$$\begin{bmatrix} 3 & 2 & 10 & 20 & 21 & 10 \\ 1 & 6 & 11 & 35 & 22 & 12 \\ 5 & 8 & 25 & 23 & 25 & 13 \\ 2 & 0 & 56 & 12 & 67 & 14 \\ 4 & 9 & 45 & 11 & 50 & 67 \\ 7 & 10 & 50 & 34 & 23 & 80 \end{bmatrix} * \begin{bmatrix} w_1 \cdot 1 & w_2 \cdot 0 & w_3 \cdot 1 \\ w_4 \cdot 1 & w_5 \cdot 0 & w_6 \cdot 1 \\ w_7 \cdot 1 & w_8 \cdot 0 & w_9 \cdot 1 \end{bmatrix} = \begin{bmatrix} \\ \\ \\ \\ \\ \end{bmatrix}$$

Check out the dimensions of the following convolutional operation

$$\begin{bmatrix} x & x & x & x & x & x \\ x & x & x & x & x & x \\ x & x & x & x & x & x \\ x & x & x & x & x & x \\ x & x & x & x & x & x \\ x & x & x & x & x & x \end{bmatrix} * \begin{bmatrix} w \cdot 1 & w \cdot 0 & w \cdot 1 \\ w \cdot 1 & w \cdot 0 & w \cdot 1 \\ w \cdot 1 & w \cdot 0 & w \cdot 1 \end{bmatrix} = \begin{bmatrix} y & y & y & y \\ y & y & y & y \\ y & y & y & y \\ y & y & y & y \end{bmatrix}$$

The result shrinks to a $n - m + 1$ dimension matrix

(6,6)

(3,3)

(4,4)

(n,n)


(m,m)

(n - m + 1, n - m + 1)


There are several setbacks of the convolution you would like to avoid

- **The result matrix is smaller**
- **The padding takes into account values in the center more than values in the edge**

Values in the center of the convoluted matrix are overused




We can augment the dimension of the matrix padding with new columns




(8,8)

*




(3,3)

=



(6,6)


You've managed to keep the dimensions our convoluted matrix



- **Same padding:** When we augment the 5x5x1 image into a 6x6x1 image and then apply the 3x3x1 kernel over it, we find that the convolved matrix turns out to be of dimensions 5x5x1.
- **Valid padding:** if we perform the same operation without padding, the result matrix is smaller

Instead of moving kernel step-by-step, we can move it further

The chart convolutes a (5,5) matrix padded to (7,7) dimension with a (3,3) and stride=2



Matrix (n, n)


Kernel (f, f)

Padding p

Stride s

$$\left(\frac{n + 2p - f}{s} + 1, \frac{n + 2p - f}{s} + 1\right)$$


RGB pictures are represented by $(n,m,3)$ matrices



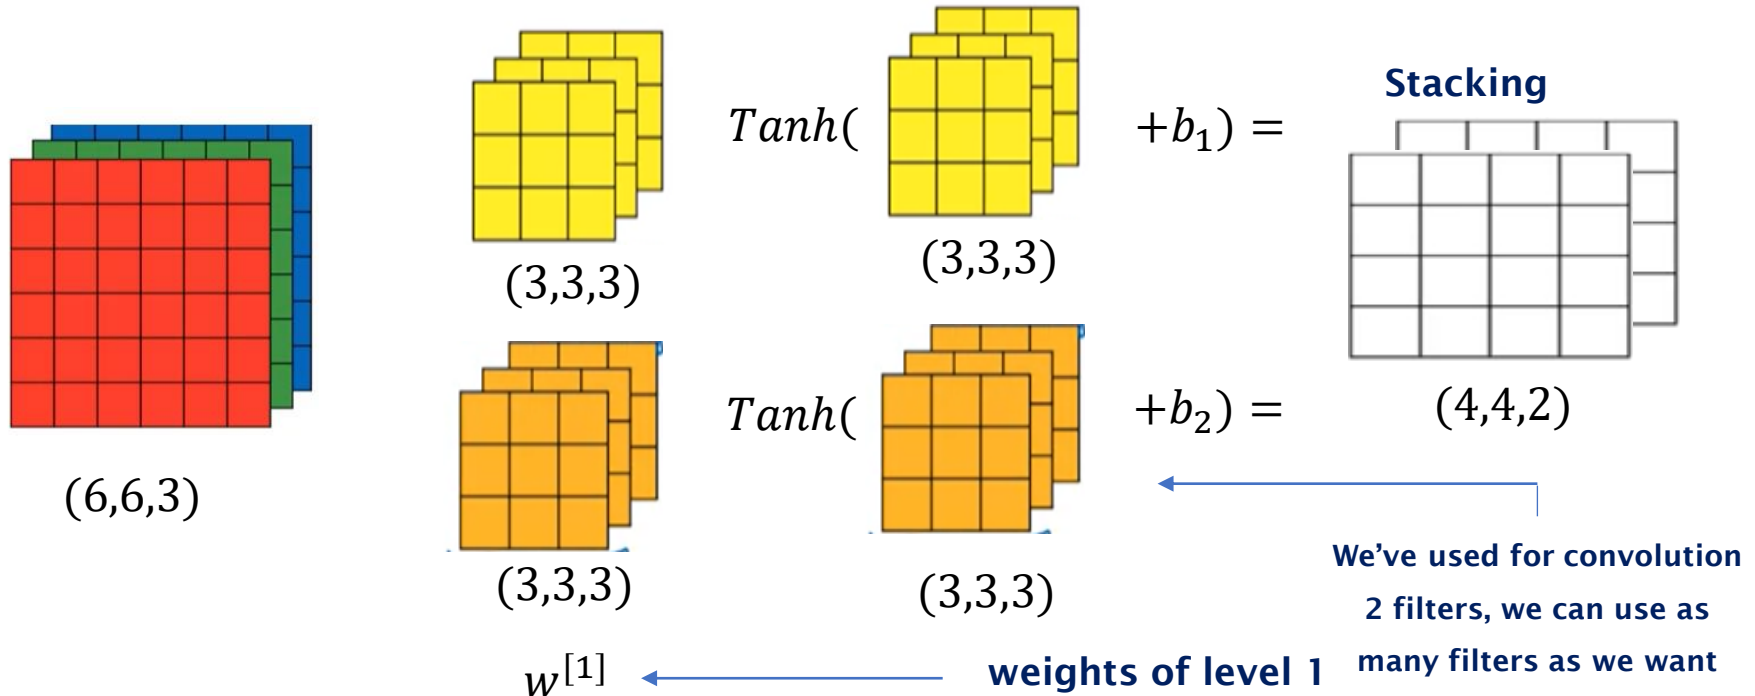
**What happened to the
third dimension?**

The first value of the convoluted Matrix will be the the sum of the first 9+9+9 values

We can apply as many filters as we want



Now we understand how Convolution Works, let's build a Convolution Neural Network






Measuring the size of your Convolutional Neural Network (CNN)

Your team decided to apply a Convolutional Neural Network with 10 filters, each filter will have a 3x3 matrix per color.

How many parameters will your CNN have?



(3,3,3)

$$3 \times 3 \times 3 = 27 + \textit{bias} = 28$$

$$28 \times 10 \textit{ filters} = 280 \textit{ parameters}$$

Small kernel size: how do we aggregate the learned micro-structures to “find” relationships at a larger scale?

Strategy in convolutional nets: stacking convolutions layers and, at the same time, downsampling the image between successive convolutions.

Strategies to downsample:

- Average pooling
- Max pooling
- Strided convolution (calculate only every n-th pixel)

Let's check the notation we're going to use

$f^{[l]}$ = filter size

$p^{[l]}$ = padding size

$s^{[l]}$ = stride size

$n_c^{[l]}$ = depth of filter


Input: $(n_h^{[l-1]}, n_w^{[l-1]}, n_c^{[l-1]})$

Output: $(n_h^{[l]}, n_w^{[l]}, n_c^{[l]})$


$$n_h^{[l]} = \left\lceil \frac{n_h^{[l-1]} + 2p^{[l]} - n_c^{[l]}}{s^{[l]}} + 1 \right\rceil$$

l = number of level, usually 3

Convolutional Neural Network Example



Useful tool to speed calculation up a make features more robust




Max Pooling: The máximo value of the área

- $f=2$ (filter size)
- $s=2$ (stride=2)

You've just applied a convolutional filter with kernel size $(2,2)$ and stride 2

Average Pooling



Average Pooling: The mean value of the área


- $f=2$ (filter size)
- $s=2$ (stride=2)

You've just applied a convolutional filter with kernel size (2,2) and stride 2

Some considerations about pooling

- Pooling doesn't involve parameters to learn, it's just a pretty straightforward technique for dimension reduction
- Padding is seldom used with Pooling
- Max pooling is much more used than average pooling

Built by Yann LeCun for digit recognition



| | Activation Shape | Activation Size | nº parameters |
|--------------------|------------------|-----------------|---------------|
| Input | (32,32,3) | 3072 | 0 |
| Conv2D 1 (f=5,s=1) | (28,28,8) | 6072 | 208 |
| Pool1 | (14,14,8) | 1568 | 0 |
| Conv2D (f=5,s=1) | (10,10,16) | 1600 | 416 |
| Pool2 | (5,5,16) | 400 | 0 |
| FC2 | (120,1) | 120 | 48120 |
| FC3 | (84,1) | 84 | 10164 |
| Softmax | (10,1) | 10 | 850 |

Total: 59758 parameters to optimize



Understanding parameters

Think about all parameters shown in the previous spreadsheet

- Figure out how all the numbers are calculated
- Why the softmax has 10 outputs (instead, for example 5)

CONVOLUTIONS

- *Understanding convolutions*
- *Building a convolutional neural network*

The `torch.nn` module provides convolutions for 1, 2, and 3 dimensions:

- `nn.Conv1d` for time series,
- `nn.Conv2d` for images,
- and `nn.Conv3d` for volumes or videos.

For our CIFAR-10 data, we'll resort to `nn.Conv2d`.

Need to provide (at a minimum):

- Number of input features: in this case, 3 for the RGB channels
- Number of output features: arbitrary, more output features, higher capacity of the model
- Size of the kernel

Please, read the documentation

CONV2D

```
CLASS torch.nn.Conv2d(in_channels, out_channels, kernel_size, stride=1, padding=0, dilation=1,  
groups=1, bias=True, padding_mode='zeros', device=None, dtype=None) [SOURCE]
```

Applies a 2D convolution over an input signal composed of several input planes.

In the simplest case, the output value of the layer with input size (N, C_{in}, H, W) and output $(N, C_{out}, H_{out}, W_{out})$ can be precisely described as:

$$\text{out}(N_i, C_{out_j}) = \text{bias}(C_{out_j}) + \sum_{k=0}^{C_{in}-1} \text{weight}(C_{out_j}, k) \star \text{input}(N_i, k)$$

where \star is the valid 2D [cross-correlation](#) operator, N is a batch size, C denotes a number of channels, H is a height of input planes in pixels, and W is width in pixels.

This module supports [TensorFloat32](#).

On certain ROCm devices, when using float16 inputs this module will use [different precision](#) for backward.

- `stride` controls the stride for the cross-correlation, a single number or a tuple.
- `padding` controls the amount of padding applied to the input. It can be either a string `{'valid', 'same'}` or a tuple of ints giving the amount of implicit padding applied on both sides.
- `dilation` controls the spacing between the kernel points; also known as the *à trous* algorithm. It is harder to describe, but this [link](#) has a nice visualization of what `dilation` does.
- `groups` controls the connections between inputs and outputs. `in_channels` and `out_channels` must both be divisible by `groups`. For example,
 - At `groups=1`, all inputs are convolved to all outputs.
 - At `groups=2`, the operation becomes equivalent to having two conv layers side by side, each seeing half the input channels and producing half the output channels, and both subsequently concatenated.
 - At `groups=in_channels`, each input channel is convolved with its own set of filters (of size $\frac{\text{out_channels}}{\text{in_channels}}$).

Quick reminder of how to load cifar10

```
label_map = {0: 0, 2: 1, 1: 2}

cifar10 = datasets.CIFAR10(
    data_path,
    train=True,
    download=True,
    transform=transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize((0.4915, 0.4823, 0.4468),
                              (0.2470, 0.2435, 0.2616))
    ]),
    target_transform=lambda label: label_map[label] if label in label_map else label
)

cifar10_val = datasets.CIFAR10(
    data_path,
    train=False,
    download=True,
    transform=transforms.Compose([
        transforms.ToTensor(),
        transforms.Normalize((0.4915, 0.4823, 0.4468),
                              (0.2470, 0.2435, 0.2616))
    ]),
    target_transform=lambda label: label_map[label] if label in label_map else label
)

birds_aeroplanes_train = [index for index, sample in enumerate(cifar10) if sample[1] in {0, 1}]
birds_aeroplanes_val = [index for index, sample in enumerate(cifar10_val) if
                        sample[1] in {0, 1}]

cifar2 = torch.utils.data.Subset(cifar10, birds_aeroplanes_train)
cifar2_val = torch.utils.data.Subset(cifar10_val, birds_aeroplanes_val)
```

← Loading pictures, converting to Tensors

← Normalizing

← Mapping labels

← Choosing only airplanes and birds

← Building Subsets

Convoluting a sample

```
# Defining de convolution operation
conv = nn.Conv2d(in_channels = 3,
                  out_channels = 16,
                  kernel_size = 3,
                  stride       = 1,
                  padding      = 0
                  )

# Converting to a 4-dimension Tensor
img_unsqueezed = img.unsqueeze(0)
img_output = conv(img_unsqueezed)

print("Size of tensor previous to the convolution: {0:s}"
      .format(str(img_unsqueezed.size())))
print("Size of tensor after the convolution: {0:s}"
      .format(str(img_output.size())))
```

```
Size of tensor previous to the convolution: torch.Size([1, 3, 32, 32])
```

```
Size of tensor after the convolution: torch.Size([1, 16, 30, 30])
```



Calculating output dimensions...

$$n_h = \left(\frac{n + 2p - f}{s} + 1 \right) = \frac{32 + 0 - 3}{1} + 1 = 30$$

Picture dimensión has been
reduced from (32,32) to (30,30)

Convolution output

Convolution generates 16 blurred pictures according to the 16 filters used.



```
for n in range(0,15):  
    plt.imshow(img_output[0,n,:,:].detach())  
    plt.show()
```



Why are all the outputs different?

Are we using the same kernel?

CONVOLUTIONS

- *Understanding convolutions*
- *Building a convolutional neural network*

Blurring: a kernel with constant weights

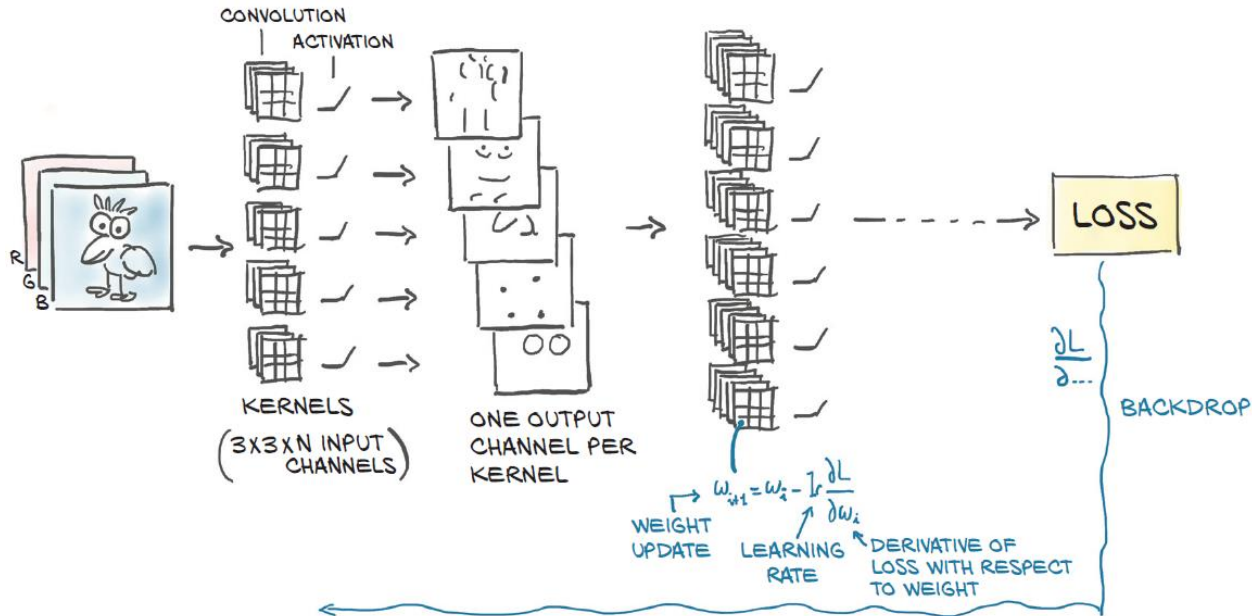
```
conv = nn.Conv2d(3, 1, kernel_size=3, padding=1)
with torch.no_grad():
    conv.weight[:] = torch.tensor([[1.0, 1.0, 1.0],
                                    [1.0, 1.0, 1.0],
                                    [1.0, 1.0, 1.0]]) / 9
    conv.bias.zero_()
```

Edge detection: differences between neighboring pixels

```
conv = nn.Conv2d(3, 1, kernel_size=3, padding=1)
with torch.no_grad():
    conv.weight[:] = torch.tensor([[ -1.0, 0.0, 1.0],
                                    [ -1.0, 0.0, 1.0],
                                    [ -1.0, 0.0, 1.0]])
    conv.bias.zero_()
```

The job of a convolutional neural network is to estimate the kernel of a set of filter banks in successive layers

- will transform a multichannel image into another multichannel image, where different channels correspond to different features



Output tensors from a convolution layer (usually followed by an activation just like any other linear layer) tend to have high values where certain features corresponding to the estimated kernel are detected (such as vertical lines).

- Through training, kernel weights adjust to detect important features
- These features, where detected, produce high value in some parts of the tensor
- With downsampling, these parts “survive” at the expense of weaker responses

Max pooling provided in Pytorch by the nn.MaxPool2d module

- Takes as input the size of the neighborhood over which to operate the pooling operation.
- Example, to downsample our image by half, use a size of 2

```
pool = nn.MaxPool2d(2)


img_output_pool = pool(img_unsqueezed)
print("Size of tensor after the pooling: {0:s}"
      .format(str(img_output_pool.size())))
```

```
IPdb [6]: img_unsqueezed.size()
Out [6]: torch.Size([1, 3, 32, 32])

IPdb [7]: img_output_pool.size()
Out [7]: torch.Size([1, 3, 16, 16])
```

Combining convolutions and downsampling

Effect of stacking convolutions and downsampling: a large cross is highlighted using two small, cross-shaped kernels and max pooling.




First set of kernels operates on small neighborhoods on first order, low-level features

Second set of kernels effectively operates on wider neighborhoods, producing features that are compositions of the previous features.

Procedure:

- Convert the 8-channel 8×8 image into a 1D vector (plus batch dimension); and
- Complete our network with a set of fully connected layers



CONVOLUTIONS

- *Understanding convolutions*
- *Building a convolutional neural network*
- *Creating custom `nn.Module` subclasses*

Subclass nn.Module

Define a forward function that takes the inputs to the module and returns the output

- Typically, forward will use other modules—premade like convolutions or customized.
- To include these *submodules*, we define them in the constructor `__init__`
- Need to call `super().__init__()`

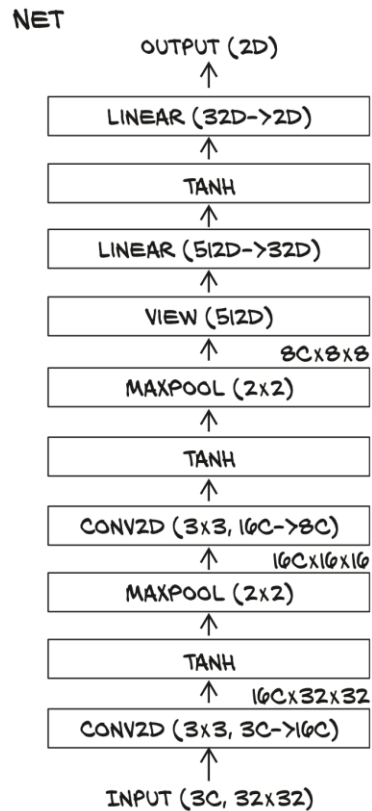
```
class Net(nn.Module):
    def __init__(self):
        super().__init__()
        self.conv1 = nn.Conv2d(3, 16, kernel_size=(3, 3), padding=1)
        self.act1 = nn.Tanh()
        self.pool1 = nn.MaxPool2d(2)
        self.conv2 = nn.Conv2d(16, 8, kernel_size=(3, 3), padding=1)
        self.act2 = nn.Tanh()
        self.pool2 = nn.MaxPool2d(2)
        self.fcl = nn.Linear(8 * 8 * 8, 32)
        self.act3 = nn.Tanh()
        self.fc2 = nn.Linear(32, 2)

    def forward(self, x):
        out = self.pool1(self.act1(self.conv1(x)))
        out = self.pool2(self.act2(self.conv2(out)))
        out = out.view(-1, 8 * 8 * 8)  # <1>
        out = self.act3(self.fcl(out))
        out = self.fc2(out)
        return out
```

Missing conversion: we need to modify the output of the second MaxPool2d using view.

We will now see a new, more versatile way of building a neural nets.

Baseline convolutional network architecture



Loading the dataset (the loader with handle batches and shuffling):

```
train_loader = DataLoader(cifar2,  
                           batch_size=64,  
                           shuffle=True)  
val_loader = DataLoader(cifar2_val,  
                        batch_size=64,  
                        shuffle=False)
```

Activating GPU

```
device = (torch.device('cuda') if torch.cuda.is_available()  
          else torch.device('cpu'))  
print(f"Training on device {device}.")
```

Training loop

```
def training_loop(n_epochs: int,
                  optimizer: torch.optim,
                  model: nn.Module,
                  loss_fn: nn.Module,
                  train_loader: DataLoader,
                  device: torch.device):

    for epoch in range(1, n_epochs + 1):
        loss_train = 0.0
        for imgs, labels in train_loader:
            imgs = imgs.to(device=device) # <1>
            labels = labels.to(device=device)
            outputs = model(imgs)
            loss = loss_fn(outputs, labels)

            optimizer.zero_grad()
            loss.backward()
            optimizer.step()

        loss_train += loss.item()

    if epoch == 1 or epoch % 10 == 0:
        print('{} Epoch {}, Training loss {}'.format(
            datetime.datetime.now(), epoch,
            loss_train / len(train_loader)))
```

← Loop

← Loading images and labels to GPU

← Forward Pass

← Calculating Loss

← Backward Pass

← Repetition

```
def validate(model: nn.Module,
             train_loader: DataLoader,
             val_loader: DataLoader,
             device: torch.device):
    accdict = {}
    for name, loader in [("train", train_loader), ("val", val_loader)]:
        correct = 0
        total = 0

        with torch.no_grad():
            for imgs, labels in loader:
                imgs = imgs.to(device=device)
                labels = labels.to(device=device)
                outputs = model(imgs)
                _, predicted = torch.max(outputs, dim=1) # <1>
                total += labels.shape[0]
                correct += int((predicted == labels).sum())

    print("Accuracy {}: {:.2f}".format(name, correct / total))
    accdict[name] = correct / total
    return accdict
```

Checking accuracy from both
train and validation datasets

Sending the batch of pictures
and labels to device

Predicting

Choosing the max value of
prediction

torch.max returns the maximum value of the given set of values and “optionally” the index of the maximum value

TORCH.MAX

`torch.max(input) → Tensor`

Returns the maximum value of all elements in the `input` tensor.

• WARNING

This function produces deterministic (sub)gradients unlike `max(dim=0)`

Parameters:

input (*Tensor*) – the input tensor.

Example:

```
>>> a = torch.randn(1, 3)
>>> a
tensor([[ 0.6763,  0.7445, -2.2369]])
>>> torch.max(a)
tensor(0.7445)
```

`torch.max(input, dim, keepdim=False, *, out=None)`

Returns a namedtuple (`values`, `indices`) where `values` is the maximum value of each row of the `input` tensor in the given dimension `dim`. And `indices` is the index location of each maximum value found (argmax).

If `keepdim` is `True`, the output tensors are of the same size as `input` except in the dimension `dim` where they are of size 1. Otherwise, `dim` is squeezed (see `torch.squeeze()`), resulting in the output tensors having 1 fewer dimension than `input`.

← Returns a tuple

The output of the model

```
IPdb [20]: outputs
Out [20]:
tensor([[ 1.3349, -1.9178],
        [ 1.9040, -1.6632],
        [ 1.5620, -1.0870],
        [ 2.0129, -1.5942],
        [-0.3770,  0.6652],
        [ 0.1299,  0.5198],
        [ 0.5858, -0.4821],
        [-1.7611,  1.8181],
        [ 3.7983, -3.1361],
        [-2.4987,  2.0394],
        [-2.5485,  2.5286],
        [ 3.0376, -2.5142],
        [ 0.2787, -0.1730],
        [-3.0452,  2.9423],
        [ 0.6357, -0.9466],
        [ 2.4407, -2.3764],
        [-3.2843,  3.0517],
        [ 0.2308, -0.1195],
        [ 3.7184, -3.4098],
        [-3.1932,  3.6036],
        [-2.5922,  1.8526],
        [ 0.3395, -0.1361],
        [ 2.0941, -1.5432],
        [-3.3572,  3.3926],
        [-2.3391,  2.0032],
```

The output of torch.max

```
IPdb [25]: torch.max(outputs, dim=1)
Out [25]:
torch.return_types.max(
  values=tensor([1.3349, 1.9040, 1.5620, 2.0129, 0.6652, 0.5198, 0.5858, 1.8181,
                3.7983, 2.0394, 2.5286, 3.0376, 0.2787, 2.9423, 0.6357, 2.4407,
                3.0517, 0.2308, 3.7184, 3.6036, 1.8526, 0.3395, 2.0941, 3.3926,
                2.0032, 2.6029, 0.2059, 1.9168, 1.4595, 3.1268, 4.0068, 3.8308,
                2.3819, 2.9784, 2.3783, 4.2831, 1.9990, 0.6077, 1.4134, 1.8785,
                2.3575, 0.1935, 2.5324, 2.9165, 1.8345, 2.1046, 1.1047, 1.5087,
                1.6084, 1.6869, 2.8519, 0.7044, 1.6892, 3.5769, 1.1302, 0.8916,
                1.9175, 0.6086, 1.4908, 0.5524, 1.8217, 4.0356, 2.7264, 3.0865],
                device='cuda:0'),
  indices=tensor([0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 1, 1, 0,
                 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1,
                 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1],
                 device='cuda:0'))
```

We can compare the position of our maximum value with the ground truth.

Remember! {0:airplane, 1:bird}

```
correct += int((predicted == labels).sum())
```

[illegible]